CURVE TURNING ANGLE PREDICTION OF CAR STEERING USING COMPUTER VISION

Ms.K.Bhavani
Assistant professor
Computer science and engineering
Kalasalingam Academy Of Research and Education
Virudhnagar,Tamilnadu
K.bhavani@klu.ac.in

M. Pavan Charan
Computer science and engineering
Kalasalingam Academy Of Research and Education
Virudhnagar,Tamilnadu.
9920004387@klu.ac.in

N. Dinath Reddy
Computer science and engineering
Kalasalingam Academy Of Research and Education
Virudhnagar,Tamilnadu
9920004447@klu.ac.in

CH. Sai Hemanth
Computer science and engineering
Kalasalingam Academy Of Research and Education
Virudhnagar,Tamilnadu
9920004383@klu.ac.in

S. Sreenivasa Rao
Computer science and engineering
Kalasalingam Academy Of Research and Education
Virudhnagar,Tamilnadu
9920004390@klu.ac.in

Abstract—Over the past few years, the number of self-driving cars has increased significantly. A dataset that includes an album of photos with the steering angle has been made available by Udacity. while operating a vehicle. Udacity's challenge sought to estimate steering angle using just the photos that are presented. To produce excellent results, We develop different models, utilizing various deep learning approaches and algorithms, to estimate steering angles based on photos.

***Keywords*—**Self-driving cars,Dataset,Steering angle, Transfer Learning, Deep learning, CNN (Convolutional Neural Network).

I.INTRODUCTION

The next ten years will see a significant economic impact from self-driving cars. In the wake of the challenge, Udacity published an assortment of pictures..Captured while operating a vehicle and the matching steering angle for a training set and supplementary sensor data (left, right, and use approximated angles based on camera angle to centre cameras). Finding a model that will minimize the error mean square (RMSE) between the angle of steering that that an actual driver really delivers, and the model's predicted outcome and prognosis was the problem at hand. The model had to be compared to a photograph obtained while operating.

The goal of the research is to develop a system that learns to drive by watching, doing away with the requirement for manually coding rules. One crucial component of the end-to-end strategy for autonomous vehicles involves steering degree prediction, which would enable us to fully utilize neural networks' capabilities. For instance, deep neural networks are able to automatically pull characteristics that assist locate the road in order to make estimations, using only the angle at which to steer as a learning signal.

Using this model, Examine the methods for predicting steering angle using temporal information. Temporal information is used by 3D convolutional layers. The subsequent model Transfer learning, This uses a superior model developed on a different data (which employs the lowest levels of an already trained models). Transferring knowledge assists in decreasing the quantity of data used for training needed to provide an excellent example.

II. Related Work

The Autonomously Ground Automobile in an Artificial Neural Network (ALVINN) systems was created by Pomerleau (1989) [13], and it was the first neural network-based self-driving navigational system. By today's standards, the model's structure was fairly basic, consisting only of a fully interconnected system. The network predicted actions by applying pixels information to simple operating situations with only a minimal barrier. Nonetheless, it showed how artificial neural networks may be used to provide complete autonomy in navigation.

An approach similar to this was published by NVIDIA in a paper last year, which benefited from ALVINN. In the publication , the authors extracted features from the driving frames using a rather simple CNN architecture. Figure 1 shows how the architecture is laid out. The extension of the information gathered was determined to be significant. The writers employed arti-generated rotations and shifts in the training set. Both to the left and right Additionally, cameras that had their steering angles interpolated were included. This system performed admirably in amazingly simple everyday situations, such following roads on a motorway and traversing flat, clear terrain.

Recently, there has been an increased effort to tackle the problems of object detection [17], image interpretation [4], and video categorization [9] by utilizing deep neural networks and RNNs as well applications of ever-more intricate CNN structures in automated vehicles. The book Collecting Spatiotemporal Knowledge Using "3D Convolutional Networks" explains how to use 3D CNN to build spatial characteristics in a collection of images or videos [18]. One of both methods described in In addition to Brief "A sample:

Deep Networking for Videos Categorization" makes use of LSTM to manage videos.

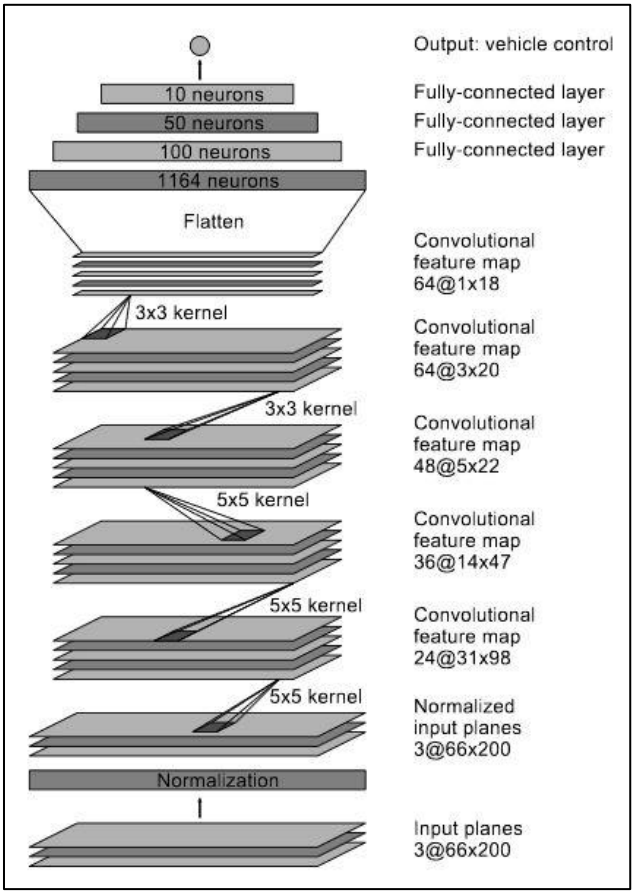

Figure 1 shows the CNN architecture from [1]. There are over 250,000 parameters and 27 million connections in the network.

The methods for creating residual connections between various layers that facilitate deep neural network training are described in "Deep Spectral Learning for Images Detection" [6] as well as "Firmly Coupled Convolutional neural networks Networks" [7].

In addition to CNN and/or RNN approaches, several research projects are using deep learning methods to address autonomous driving problems.. Managing self-navigating cars is a different job since it involves video predictions. One method recommended by Comma.ai [14] for learning is to use an automobile simulation. With a technique that combines a Generative Adversarial Network (GAN) and a Variational Auto-encoder, (VAE) [11] [5] Using prior frames as a basis, their method may continue to forecast realistic-looking video for several frames. even though the transition model was tuned for free operate within the pixel domain.

Additionally, autonomous driving has also made use of deep reinforcement learning (RL) [3], [15]. Until recently, RL has not proved successful for automotive applications. However, new research indicates that deep learning algorithms can learn accurate depictions of the surroundings. It was shown by picking up games like Go and Atari by DeepMind by Google [12], [16]. Motivated by these efforts, [3] has put forth a framework for driverless vehicles that uses rich RL. They can incorporate RNN into their structure. for information integration, which permits the vehicle to manage situations that are only partially observable. By using the peek and response network to guide its CNN kernel toward the provided information places that are relevant to the operating activity, the architecture also incorporates attention models.

III. PROPOSED MODEL

3. Methods

We developed a pair of algorithms. While the second one employs LSTM recurrent layers, the initial one uses three-dimensional layers of convolution. Two-dimensional transfer learning is used in the second model. a convolutional layer-based model that has been trained in which layers are not permitted during the initial training.

3.1. 3D Convolutional Model with Recurrent LSTM Layers and Residual Connections

3.1.1 3D Convolutional Layer

The only way a 3D convolutional layer differs from a 2D convolutional layer is that it now has a third dimension, depth (temporal), in addition to height and width. Rather of use a 2D filter (assuming we disregard the channel dimension temporarily) shifting inside the picture Now, in addition to depth, width, and height, we have a three-dimensional filter that travels across both width and height. When an input is in the form D1, H1, W1, C, the output would be in the form D2, H2, W2, F, F represents the quantity of filters. Calculations of D, H, and W2 can be made using the stride and cushioning lengths.

3.1.2 Residual Connection

The residual connections are used during training to support deeper neural networks, which could struggle to transfer the changing gradient throughout every level. Utilizing network levels is the concept behind residual connections to support an ongoing mapping. Rather than explicitly trying to suit a particular fundamental mapping. Lack of persistent connection: Plot explicitly H(G(x)) With an ongoing partnership:

Suggested $H(G(x)) = F(x) - x - G(x)$

In essence, the changing gradient might pass throughout various levels. Skipping reduces the gradient's effective number of layers. must go through in order to return the entire distance. When skipping in this way, issues with the backpropagation process in really deep networks like the issue with fading gradients.

3.1.3 Spatial Batch Normalization

Many of the issues related to correctly initializing neural networks are eliminated when activations throughout a network are purposefully made to adopt a unit gaussian distributions from the beginning of learning, as is the case with batches of normalization (see [8]). Spatial batch normalization not only normalizes across samples but also along the dimension axis on pictures. Applying this spatially normalization batch comes after each 3D convoluted level.

3.1.4 Long Short-Term Memory with Recurrent Neural Networks

Loops in RNN enable the persistence of information. Using prior video frames as a basis, this might be utilized to comprehend the current frame. One issue with vanilla RNN is that because of the gradient disappearing, it is unable to

train to connect data that is too far. The unique cell state of the LSTM functions as a conveyor bell facilitating the free flow of information with minimal interactions. Visual features are extracted using 3D convolutional layers, and then the sequential relation is captured by feeding the extracted features into LSTM layers.

3.1.5 New Architecture

Production systems for autonomous automobiles may find it useful to incorporate temporal information. Knowing the details of the prior frames, for instance, would be helpful if the camera sensor was fully saturated while it was facing the sun. enable a far more accurate forecast than relying on the steering prediction of angle limited to the saturated frame. As mentioned Previous methods employ recurrent and 3D convolutional layers of data to correlate temporal information. We used both of these approaches to applying chronological data in our model. We employed when building this model, the concept of residual connection [6]. More of the gradient is able to travel via these links. via network by merging several layers. The instance comprised five video frame sequences that were each shifted by The input consists of one frame (5x5x120x320x3). The principles were.

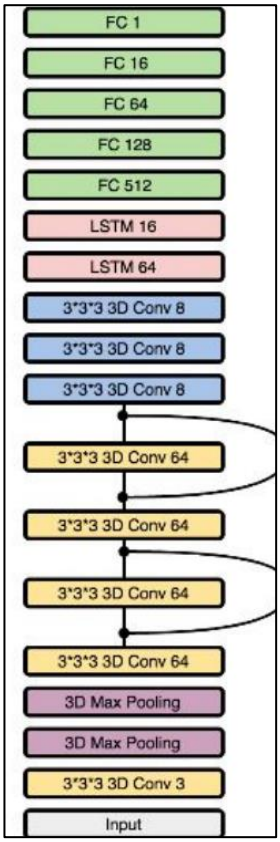

Figure 2. 3D convolution system featuring recurring layers of LSTM and remaining connections

3.2. Transfer Learning

The concept of transfer learning is based on the notion that features acquired at lower levels of the model are probably applicable to another dataset; for instance, lower-level features like edges could prove helpful for the fresh data set. This allows us to use excellent models that were developed on already present big data sets. One of the previously trained models, ResNet50, performed well on this dataset; It was trained using the first fifteen ResNet blocks and ImageNet. weights (i.e., the first 45 individual layers out of 175) were prevented from updating. A stack of completely linked layers of 512, 256, 64, and 1 distinct unit, respectively, was connected to the output of ResNet50.

Learning only the last five blocks produced poor results that hardly improved when every input was expected to have a steering degree of zero. Additionally, learning every layer generated less favorable outcomes on the validation set than when the first 45 (after 32 epochs, or 0.0870 on the validation set). Images of 224x224x3 that had been slightly stretched a downsized from the original Udacity data were used as input by the model. Mirrored images were the only augmentation offered for this model. Owing to the input's size limitations, Cropping was not employed in ResNet50 because it required stretching. the picture. The pretrained model's filters may not perform as well on the extended data (RMSE of 0.0891 on after 32 epochs; the validation set after) since they were not trained on stretched images. Additionally, employing the left and right cams from the set used for training proved to be insufficient for the 32 training epochs (0.17 RMSE on the validation set).

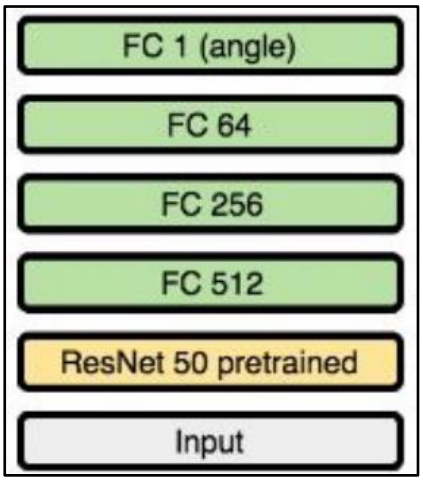

Figure 3. The transfer learning model's architecture.

The data set we utilized was produced by NVIDIA's DAVE-2 System and supplied by Udacity[1]. In particular, the vehicle has three cameras mounted behind the windshield. vehicle for gathering information. When the steering angle is employed, simultaneous, time-stamped video is collected by the cameras. directly from the driver. Coach receives this directive for direction. To ensure that the system is not dependent on the car's geometry, they express the steering instruction as 1/r. The turning radius, r, is the measurement in meters. They employ 1/r instead of r to avoid a singularity from occurring when driving directly (the turning radii for going straight is infinite). From turning left (with no positive values) to right turns through zero, 1/r flows smoothly.

possesses 5615 frames. The image's original resolution was 640x480. Five distinct driving videos provide the training images:

1. 221 seconds, bright sunshine, numerous lighting adjustments. Smooth transitions at the start, uneven shoulder lines, concludes with a divided highway and lane merge
2. The highway is divided into two lanes, with uneven shoulder lines, a lane merge at the end, and a traffic signal that is green. It also has very tight turns with little visibility for the center camera, direct sunlight, and rapid elevation changes that result in steep gain/loss during the summit. becomes a divided highway swiftly returns to two lanes at 350 s.. 99-second split highway portion of the return journey above the pea212 seconds, two-lane road with a guardrail, shadows in Although training can be difficult at first, it usually becomes simpler in the end. The amount of traffic on a divided multi-lane motorway with a fair pace is 371 seconds Typical photos under various lighting, traffic, and driving conditions are displayed in Figure4.

Figure 4. Sample pictures from the collection. From left to right, there will be strong sunlight, shadows, a sudden left turn, a straight uphill stretch, and lots of traffic.

IV.METHODOLOGY

4.1. Data Augmentation Methods

4.1.1 Enhancement of Brightness

sporadic luminosity adjustments are made to mimic various lighting scenarios. We first convert images to HSV, then scale them up or down to create augmented images with varying brightness. returning to the RGB channel after going down the V channel. Here are some examples of enhanced pictures.

4.1.2 Enhancement of Shadows

Random shadowing casts across images. The aim is that even in the event that the camera is hidden (by dust or rain, for example), the model should still be able to forecast the proper steering angle. Random points are chosen in order to apply this.

and making all of the points on a single aspect of the image darker.

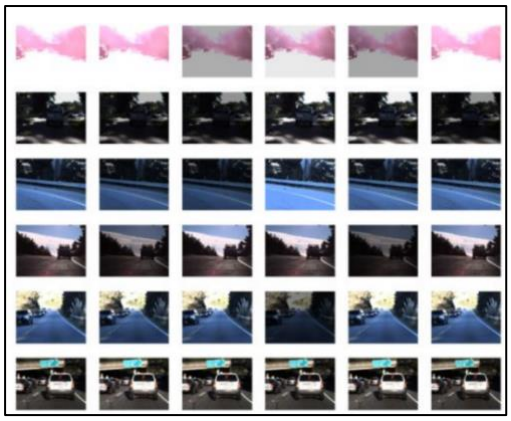

Figure 5. Examples of brightness augmentation

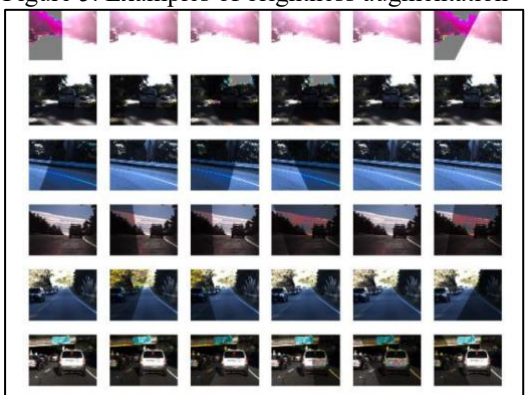

Figure 6. Shadow augmentation examples

4.2. Data Augmentation

We tested various methods of data augmentation using the architecture from NVIDIA [1] to create a baseline for our study. The Udacity challenge teams mentioned additional data augmentation, in addition to the initial NVIDIA scientists. Understanding It would be helpful to know which augmentation techniques are the fastest and most computationally accessible for training our new models. The earlier NVIDIA design is seen in the. The input dimensions for this model were 120x320x3, and its batch size was 32. The NVIDIA release did not explain how the loss of function was improved. For this study, we used the Adam preset configurations.

Three distinct augmentation levels were looked at. The initial image was just slightly altered, with the top portion of the image cropped and some accidental flips. The bias towards the right is eliminated by randomly flipping the input photos. appears in the data. resizing the image's top removes the sky, which has no place in the picture. part in predicting the steering angle and how to turn. An additional version of enhancement was the same as the minimal amount version in addition to minor image brightness adjustments, shifts (25 pixels), and rotations (-5 to +5 degrees). More dramatic effects from the second version were used in the final, more intense version of the augment, such as big shifts, large shadows, and large angle rotations (up to 30 degrees). Changes in brightness were applied. Findings from this trial.

The root mean square error (RMSE) for the validation set with 32 epochs of data augmentation using the NVIDIA architecture.

When extensive data augmentation was applied, an orientation angle of 0 for all the information was predicted with extremely poor results. The smallest enlargement yielded the best results, but the moderate augmentation yielded satisfactory results as well. The only explanation for these results would be to need thirty-two training epochs. A heavy augment could be challenging. for the model to detect such significant changes. Likewise, Perhaps the moderate version fared better than the minimal variation over several eras. A section later, a visualization of These assessments will be looked at. In selecting our latest models, we have to make only small augmentations.

4.3. Training Process

4.3.1 Loss Function and Optimization

Each model that was used was subjected to the mean-square-loss function. Regression problems often use this loss function. A large deviation is greatly affected by the MSE. Equation 1 demonstrates that the fundamental form of The sum of the squared differences between the actual and expected values is what this function represents. The Udacity is challenge scores were made available to the public. as the MSE squared, or more simply, as the root mean square errors (RMSE).

$$MSE = \frac{1}{n}\sum(y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n}\sum(y_i - \hat{y}_i)^2}$$

The Adam optimizer [10] was employed to maximize this loss. For deep learning applications, this optimizer is frequently the preferred option. Typically, this optimization technique significantly surpasses more general stochastic gradients decent techniques in terms of functionality. Initial testing of these models demonstrates that the losses The level change lowered after a few epochs.

Even so, Adam uses the decay of the learning rate to calculate an adaptive learning rate. Deterioration rate of the From 0 to the divided learning rate, the optimizer was updated. via the quantity of epochs. The remaining default settings for the During training, the Keras Adam optimizer performed well. (learning rate = 0.001, $\beta = 0.9$, $\beta 2 = 0.999$, $\sim = 1e - 8$) decay is equal to batch size / learning rate.

4.4. Feature Visualization

Saliency maps let us examine the elements of a picture that our networks think are significant. These maps can show how the gradient reappears in the image, emphasizing the areas that stand out the most. In a recent paper, NVIDIA employed a similar tactic [2].

4.4.1 Data Augmentation Experiment

what these models considered to be significant. The lane markers were deemed significant by the minimum model. More lane markers were determined to be significant in the moderate model; yet, This model appeared to have noisier prominence maps, which could account for its marginal reduced efficiency. In the heavy model, hardly any regions were determined to be noteworthy, which makes sense given its subpar execution.

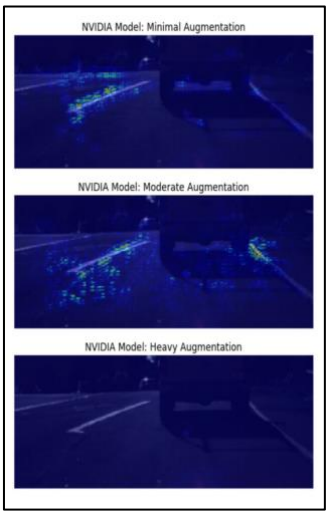

figure: data augmentation example

4.4.2 3D Convolutional LSTM Model

Convolutional neural networks (CNNs) have emerged as a powerful tool for predicting steering angles in self-driving cars. CNNs are adept at extracting spatial features from images, making them well-suited for analyzing the visual information captured by a car's front-facing camera. By feeding a sequence of images to a CNN, the network can learn to identify relevant
features and patterns that inform steering angle predictions.

Looking at one of the footage clips that were introduced into the simulation, we can observe that in Figure 10, the conspicuous features vary from frame to frame. This implies that there are substantial changes between the frames. It seems as though the key elements shift from pixel to pixel.

Illustrates how this series of frames can be collapsed into a single image. This is easier to see in the compacted version. The emphasized key characteristics do gather near road markings, but they also congregate near the shadows cast by other cars. This model might make use of details about the vehicle ahead to enable steering angle to forecast in addition to the road signs.

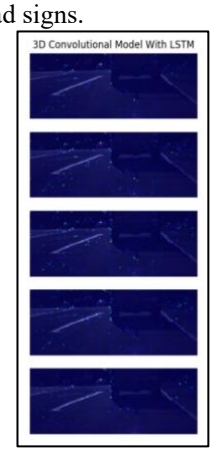

Figure: example of convolution lstm

V.RESULTS AND DISCUSSION

The outcomes of our study reveal promising advancements in the implementation of our vision-based steering angle prediction model. This model relies on camera analysis of the road area, enabling the autonomous steering of the vehicle within the lane boundaries. The primary focus of this research was to enhance the robustness of the system across diverse scenarios.

One of the primary challenges addressed in our study was ensuring the system's resilience in a variety of situations. This included adapting to illumination changes, navigating both curved and straight roads, managing urban and highway driving, responding to dynamic traffic conditions around the ego vehicle, and managing various artifacts on the roads, such as cracked surfaces.

In scenarios with changing lighting conditions, the model highlighted the ability to adjust and maintain accurate steering predictions. The challenges posed by different road geometries, including curved and straight sections, were met with success, highlighting the adaptability of the vision-based steering model.

VI.CONCLUSION

In conclusion, our study demonstrates the feasibility and effectiveness of a vision-based steering angle prediction model for autonomous vehicles. The successful navigation of diverse scenarios, including challenging lighting conditions, varied road geometries, urban and highway driving, dynamic traffic conditions, and road artifacts, underscores the model's robustness and adaptability.

While the results are promising, there is always room for further refinement and enhancement. Future research could focus on optimizing the model's performance under more complex and dynamic scenarios, improving real-time responsiveness, and exploring additional features to enhance overall system accuracy. Overall, this study marks a significant step forward in the development of vision-based autonomous driving systems.

VII.References

[1].D. P. Kingma and M. Welling. Auto-encoding variational bayes. arXiv preprint arXiv:1312.6114, 2013.

[2.]V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, and M. Riedmiller. Playing atari with deep reinforcement learning. arXiv preprint arXiv:1312.5602, 2013.

[3] D. A. Pomerleau. Alvinn, an autonomous land vehicle in a neural network. Technical report, Carnegie Mellon University, Computer Science Department, 1989.

[4] E. Santana and G. Hotz. Learning a driving simulator. arXiv preprint arXiv:1608.01230, 2016.

[5] S. Shalev-Shwartz, S. Shammah, and A. Shashua. Safe, multi-agent, reinforcement learning for autonomous driving. arXiv preprint arXiv:1610.03295, 2016.

[6] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, et al. Mastering the game of go with deep neural networks and tree search. Nature, 529(7587):484–489, 2016.

[7] C. Szegedy, A. Toshev, and D. Erhan. Deep neural networks for object detection. In Advances in Neural Information Processing Systems, pages 2553–2561, 2013.

[8] D. Tran, L. Bourdev, R. Fergus, L. Torresani, and M. Paluri. Learning spatiotemporal features with 3d convolutional networks. In The IEEE International Conference on Computer Vision (ICCV), December 2015.

[9] Udacity. An open source self-driving car, 2017.

[10] J. Yue-Hei Ng, M. Hausknecht, S. Vijayanarasimhan, O. Vinyals, R. Monga, and G. Toderici. Beyond short snippets: Deep networks for video classification. In Proceedings of the IEEE conference on computer vision and pattern recognition, pages 4694–4702, 2015.

[10]1.E. Reinhard, M. Ashikhmin, B. Gooch, and P. Shirley, “Color transfer between

[11]REFERENCES images, ” IEEE Computer Graphics Applications, vol. 21, no. 5, pp. 34–41, 202.

[12]S. Beigpour and J. van de Weijer, “Object recoloring based on intrinsic image

[13] F. Pitie, A. C. Kokaram, and R. Dahyot, “Automated colour grading using colour distribution transfer, ” Comput Vis Image Underst, pp. 123– 137, 2020.

[14] H. Chang, O. Fried, Y. Liu, S. DiVerdi, and A. Finkelstein, “Palettebased photo recoloring, ”ACM Transactions on Graphics (Proc. SIGGRAPH), vol. 34, no. 4, 5. 2021.

[15] M. P. Rao, A. N. Rajagopalan, and G. Seetharaman, “Harnessing motion blur to unveil splicing, ” IEEE Transactions on Information Forensics and Security, vol. 9, no. 4, pp. 583– 595, 2020